

Anchored decoding: per-prompt KL stays under budget across all  $K$

